Comvi - Comparative Visualization of Molecular Surfaces using Similarity-based Clustering

Wilhelm Buchmüller, Shoma Kaiser, Damir Ravilija, Enis Spahiu

Figure 1: Screenshot of a running comvi instance

Abstract— The goal of this paper is to show the reader the abstract methods and concrete applications that were used to extract and compare features and rank the similarity of the molecular protein maps. Further we present a new method of how the won data can be visualized on high resolution and large displays with dynamic interactions. The paper describes the process and the approaches that were taken to solve this task.

Index Terms—Clustering, Similarity, feature extraction, Visualization, high-resolution display, Powerwall,MegaMol, VISUS

◆

1 INTRODUCTION

TODO: add sections to tasks in this section Over the span of 6 months we, the authors of this paper, have researched and implemented a comparative clustering of molecular maps. The task consisted of several parts: This work was based on the MegaMol project. MegaMol is a simulation tool developed by the Universitt Stuttgart and the TU Dresden(?) **TODO:** check this fact[1]. It can be used to visualize particle data, simulations on atomic scale and other molecular processes. Due to its modular nature, it can be extended with modules to interact with other modules. In this paper we guide you through the **TODO:** verbose MSMCLUSTER plugin, its capabilities and inner workings.

The next task was to retrieve molecular image data through existing MegaMol plugins **TODO:** cite moleculr maps[7]. For this a special binary of megamol was compiled and will be released in a separate project **TODO:** make github link tocomvi public [1]. The next task was to extract a expressive feature vector from those images and to find a metric to cluster them by similarity.

The last task which was also developed in a plugin in MagaMol

Wilhelm Buchmller
buch.willi@googlemail.com
Enis Spahiu
enis.spahiu@hotmail.de

Shoma Kaiser
shoma.kaiser@googlemail.com
Damir Rwilja
st144286@stud.uni-stuttgart.de

[1] https://github.com/aiosin/comvi
Last retrieved: May 23, 2018

was the visualization on the VISUS POWERWALL. The POWERWALL is a very high definition display that can be used to visualize large data(sets). Due to its size its possible to display much more information than on a regular screen.

The POWERWALL also supports a tracker device that can transmit 6 degrees of freedom, so for the interaction step we had more freedom to work with than with traditional human interface devices (HIDs), for this last step we also researched the possible interactions with the unique tracking devices with **TODO:** phrasing the POWERWALL.

Over the course of the next few pages you will learn how we approached these challenges and how we (attempted) solved them. You will find out what worked and what didn't.

TODO: Part initial challenges and encountered problems einpflegen

It is clear that the task required from us that we learn how to compare the images, measure the distances between the images, and cluster these images. The given task required that we use a similarity based clustering algorithm.

Initially we were given the choice we could either chose to find similarities and cluster the proteins in the .pdb format or given as bitmap image generated by the MolecularMaps plugin in MegaMol **TODO:** cite megamol protein image[7].

Since handling image files which give information about the protein in two dimensions was easier than dealing with the pdb file format which results in three dimensional visualizations we decided to start with a two dimensional approach.

2 RELATED WORK

TODO: cite kolesar for clustering Clustering proteins by similarity or at least comparing individual proteins has been subject of existing work.

The paper [4] already had similar approaches to our results. Kolesar et al. used a 10 dimensional feature vector based on invariant image moments defined by hu **TODO:** cite hu image moments [2]

Another approach for 3D protein data were 3D zernicke moments explored by **TODO:** 3d surfer [8]. The approach is basically the same as in Kolesar et al. but the TODO used Zernicke moments instead of traditional image moments and extended them to three dimensions.

3 CLUSTERING

TODO: write some intrduction to clustering, aufteilen in subsections

3.1 Approaches to the Clustering-Problem

Right of the start we had several ideas of how we could approach this problem. With the recent trend in machine learning we had a couple of ideas of how we could determine a similarity metric between two images or classify an image into a more usable vector of data.

We ended up using a higher dimensional feature vector described in **TODO:** give section label to determine the similarirty between two protein maps because we didnt manage to train a custom model in the given timeframe, due to inexperience **TODO:** this can be said better and non existance of labeled data.

But our relatively spartan results with a pretrained Imagenet **TODO:** cite imagenet [6] **TODO:** cite darknet publication model let us to believe that given the knowledge on the subject and humanly labeled data (based on known featuresit **TODO:** LEFT OFF HERE shoud be definetly possible for this specified task to find a machine learning solution using neural networks/autoencoders.

3.2 Finding a feature vector to cluster the images

The challenge of finding a good feature vector was/is to find good features which are **TODO:** aussagekrfting about the image.

The following procedure after finding/determining/calculating the feature vector for a given image is to apply some sort of dimensionality reduction to project a higher dimensional vector onto a 2D or 3D plane.

This has multiuple advantages. First If the dimensionality reduction works as intented one find out after applying the dimensionality reduction if similar looking items are positioned next to each other.

Another reason is the curse of dimensionality. As we all know in higher/infinite dimensional spaces, otherwise unexpected things start to happen such as the euclidian distance or mathimatically put the

$$L_2 \tag{1}$$

-norm loses relevance since

$$\lim_{n \to \infty} x^n = 0 \quad \forall x \in [0, 1) \tag{2}$$

Simply put, otherwise very similar values get skewed to zero.

So we have to come up with other solutions to this problem discussed in **TODO:** put clustering section here

TODO: back to finding the feature vector After looking at a small subset of molecular map images from a variety of proteins we determined that we needed to extract feature information about the follwoing properties of a given image:

Color distribution, Shape, Texture and image moments

The initial idea was that the color distribution gives information about the color palette in the image, the extracted shape features should give information about the the biggest n shapes in the picture, the texture feature should differentiate between smooth and rough texture and everything in between.

The image moments were chosen as a goto approach to extract invariant features from the image which has been proved to yield results as described in [4]

The exact image features that we extract from the image are the following:

3.3 Invariant Image Moments by Hu

The set of invariant Image Moments discovered by Hu et. al **TODO:** find out if hu moments just him or others or et al. are rotation, translation, scale and trasformation **TODO:** find out if correct invariant. This allows us to determine if an image I_a is similar to another image I_B if I_B is equal I_A and simply rotated by 30 **TODO:** put the value in degrees there

The (continouus (spelling?) Image) Moments over two dimensions at their core are defined as such:

$$m_{i,j} = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x,y) dy dx$$

. When dealing with discrete values like we find them in an (RGB/GS) Image we use sums insteads of integral so we get this:

$$m_{i,j} = \sum_{i=0}^{n} \sum_{j=0}^{m} f(i,y)$$

.**TODO:** get definition of moment

Hu further defines his moments as such:

TODO: give the definition of all 7/8 moments (?) because this is going to take up a lot of space.

Kolesar ended up using these moments **TODO:** put in kolsar moments, in our case the moments **TODO:** put the hu moments here

3.4 Color palette/histogram

The goal when extracting color palette was to reduce the big color space that is present in any of the molecular maps and get a few distinct **TODO:** aussagekrfting colors from that range.

To achieve this we extract a histogram of each color channel of the RGB images. Each channel is then represented as a greyscale image. We then create for each channel a histogram of the luminance intensities with in our case 16 bins. This "reduces" the $128 * 128 = 16384$ dimensions to just 48 dimensions for our image.

Alternatively you can take another approach described here

The results of both approaches will be discussed in **TODO:** cite

3.5 Texture

TODO: better introduction After looking at a toy dataset of molecular maps we noticed that many samples had a distinct roughness that looked like they could be used to classify their texture.

We ended up using the haralick textural features. The haralick features work with a grey level correlation matrix (GCM). The gcm for a given greyscale image is defined as such:

put the formula here, this is spaceholder

Also we can take more features from that gcm than just the **TODO:** what exactly can be taken from this matrix which gives us more features to work with.

we also computed this on every channel yielding us another x features **TODO:** expand texture features.

Alternatively one can also use the tamura texture features. the tamura features are different to the haralick featues because they **TODO:** what exactly. They were primarily researched/developed/created to **TODO:** what exactly.

3.6 Shape

We did not end up using the shape features described here but, since we spent quite some time getting them to work **TODO:** rephrase I think it is important show how and why they were created.

The shape features we hoped to use were fourier descriptors[2] [3]. In short, we get the centroid-distance curve of a distinct shape and

---

[2] http://fourier.eng.hmc.edu/e161/lectures/fd/node1.html

[3] http://demonstrations.wolfram.com/FourierDescriptors/

take the fast (discrete) fourier transform of that curve. The technique has been proven to work for shape recognition in earlier works [3].

Our approach was the following. Since our molecular maps are extremely complex (image-wise) we first need to divide the image into discrete image region. We preferably want to segment our image into *n* colors. To recap: our everyday monitor does support 24 bit colors, 8 8 bit per channel. that yields us $2^{24} = 16777216$ colors. We'd be lucky to find a countour in this color mess. So our approach is to squash the color space down to a couple of colors. This proces is called color segmenatation. We use the k means algorithm to put every

Alternatively the mean-shift algorithm can be used. The algorithm is a kernel based clustering algorithm which operates sort of like the gradient descent algorithm. Each point has a weight and initally all the points are laid out on a grid (or your dimensions next best spatial representation) each value attracts with a constant force all other neighboring values so they all suck each other in discrete clusters depending on the parameters (bandwith size and kernel weight) A comparison of color segmentation of mean shift and k menas can be seen here. **TODO:** give image for comparison

But our problem of color segmenatation was stil not solved. Especially in the k means algorithm we encountered a lot of smaller "irrelevant" shapes that we didnt want. So one way of getting them out of the way was to squash the color space further down to two or three distint colors in the image.

Else we'd have to dynamically dilate and erode **TODO:** give link to dilation and erosion the image until we can detect with a matrix labeling algorithm / floodfill algorithm that we really have only a handful of shapes in the image.

For the shapes we then extract the contours and compute the centroid of that shape. The centroid is the geometric center of mass of the shape or in other areas of science known as ..... the first moment

Now we have the centroid and the countour of the shape that we want to classify we need to form the contour centroid distance curve. The contour centroid distance curve is simply a disrete list of values of the contour **TODO:** clafiry or the list of the euclidian (in two or three dimensions, why this doesnt work in higher dimensions look up **TODO:** curse of dimensionality) distances.

We then normalize the values by diving all the values by the biggest value in the array, resulting in a array that is in range(0,1]

Without loss of generality depending on whhich curve we chose we take the fast discrete fourier transform of that signal, which gives us for each sine/cosine coefficient a weight how that component contributes/weighs in the signal.

Without any proof that would go beyond the scope of this paper we can say that the gained features are rotation scale and translation invariant. Its rotation invariant because ... its scale invariant because we normalize it, its translation invariant because the distance curve is translation invariant (wihtout proof).

Another shape feature can be won, if the shape object is given as a greyscale binary matrix. From this matrix one can extract Image Moment features like the one by Hu as described in **TODO:** give section to hu moments

We did not end up using the shape features, because we could notice any improvement over the other features that we used, but it is important to note that they exist and what their capabilites are incase someone will continue our work on this topic. **TODO:** rephrase.

3.7 Finding the best performing similarity measure

After our feature processing on an image we end up with an feature matrix where the rows are samples and the columns are the features. aaaaa

$$A \in \mathbb{R}$$

3.8 Findin the best performign dimensionality reduction method

For testing purposes we used every dimensionality reduction method we could find that looked halfway promising. We tested on datasets of various sizes and content.

We tested our results on the Oxford flower dataset and bmw dataset and our large (3000 images) molecular maps dataset and a medium version of that dataset (300 images) and a small dataset for fast testing purposes.

As you will find in our comvi repository **TODO:** give link the dimensionality reduction routine consists of passing it a features array with **TODO:** beautify shape samples*features In the routine we additionally scale the data because **TODO:** why exactly ?!!??. We then return the reduced feature array back to the callee. The callee can then apply further some functions on this....

After inital testing with linear dimensionality reductions we notice that our results just wouldnt **TODO:** phrasing: converge , so we again tried any non linear dimensionality reduction technique we could find.

We tested our data with LLE, MDS, kernel PCA, ISOMAP and t-SNE. After testing with a small user group that would judge how good the clustering was done by the technique we concluded that t-SNE performed the best.

Further testing leads us to believe that t-SNE is the best dimensionality reduction procedure for this problem.

t-SNE is a randomized t bla Stochasitc neighbor embedding process. Its works by ... . But because the process requires random variables we decided to set the random seed of our program sowe can get deterministically the same results across different machines.

3.9 Clustering - Algorithms and strategies

TODO: this will be the main section for the clustering topic Clustering of data is defined as the process of labeling n distinct values into a set of subsets with the numbers of labels being less than

There are several distinct clustering algorithm types. Hierarchical clustering, densitty based ..**TODO:** finish

3.10 Finding the best performing clustering algorithm

Our approach to finding a clustering algorihtm was the same as the one that we used **TODO:** sentence and grammer !!! to find the best performing dimensionality reduction algorithm.

Except with a little modification. Kolesar et al. [4] used the k nearest neighbors algorithm over their 10-dimensional feature vector to cluster their data, as they discuss their results. We decided to further use this approach to cluster our molecular maps instead of the molecular protein tunnels which Kolesar et al. analyzed. **TODO:** phrasin.

In the early stages of the projcet we intended to use the DBSCAN clustering algoirtm. The key feature to this algorithm is that it automatically determines the optimal number of clusters that it detects in the given data.

The algorithm is a density based spatial clustering algirhtm that uses a similarity metric between two items that are to be clustered defined over the domain $[0, 1]$. So the ways of computing the similarity is given more freedom with this approach.

You could use the euclidian distance between two feature vectors or the cosine distance or feed the two images to a siamese neural network/autoencoder **TODO:** give information about siamese neural networks and autoencoders which will have learned the similarity metric between the images.

You could also usethe householder/hausdorff distancebetween two higher dimensional vectors.

We opted against DBSCAN for our primary clustering algorithm because in the early stages of the project the computation of the *nxn* similarity matrix was very computationally expensive and took quite a bit more time than the approaches that we demonstrate in the next paragraphs

There are similar approaches **TODO:** finish sentence

For our testing we used k nearest neighbors with the elbow method and a cutoff at 85 % **TODO:** what exactly

spectral clustering, chinese whispers ...

Figure 2: Early test wihh the Oxford flower dataset **TODO:** cite oxford flower dataset


Figure 4: final test with the tsne dimensionality reduciton and the mean shift clustering algorithm


Figure 3: Early test with the bmw car dataset **TODO:** cite stanford car dataset[5]

Figure 5: final test with the tsne dimensionality reduciton and the mean shift clustering algorithm

We ended up with the mean shift algorithm. It also is a density based algorithm and much like the DBSCAN algorithm is also determines the optimal number of clusters for the given parameters but it operates with a kernel sampled on each point in the plot and calculates the center of mass so to speak of the points clouds, for a animated demonstration lookt at **TODO:** give link to mean shift visualization. For a demonstration of the mean shift algorithm look at Figure 4.

3.11 Testing the feature vector with other datasets

4 VISUALIZATION AND INTERACTION

TODO: write some introduction about powerwall, trackingsystem, visus etc.

4.1 Approaches to the Visualization-Task

Our aproaches to visualizing the given clusters were the following, the reader is reminded that we are not just visualizing the clusters on a "normal machine" but rather the POWERWALL, a projected display with effectively 6-24 times the resolution of a consumer grade display. Details on the POWERWALL can be found **TODO:** cite powerwall publicatoin here if available. [9]

If we are given so much pixel real estate we are given the freedom to draw smaller pixels since we still will be able to see them on the POWERWALL.

During the duration of the project the idea of a 3D visualization was discussed among the team, but we settled for a 2D visualization. This had a couple of reasons.

While the interactions with the data in 3D would have been more fun since we had more degrees of freedom to work with. But we could not find a way to present the data in a way such that with just

a glance the user could intuitively interpret the data that would be displayed on the screen.

... So we decided to settle for a 2D visualization. Our approach is rather boring but it works. On startup we display nothing, if the user chooses his supplied image data and the algorithms used to cluster the pictures he gets the option to start the visualization.

The visualizatoin consists of displaying the image of a cluster representative with a simple rectangle. This way the user knows exactly what to expect to be in the cluster.

If the user wishes to have a closer look at the images in the given cluster he can click onto the representative and will get a view of all the images in that cluster.

We decided after testing with toy test and real datasets that one level of subclustering is enough, after 2 levels of subclustering the clusters get **TODO:** find better expression noisy and ambiguous.

In both the main and subcluster view the representatives are visualized with "force directed layout", to avoid

4.2 Approaches to the Interaction with and mouse and the Powerwall tracking stick (?) **TODO: cite correct pub and use correct name**

TODO: @shoma @enis @damir interaktionsmglichkeiten schreiben **TODO:** rephrase The tracking system that was provided to us from the VISUS of the University of Stuttgart was the **TODO:** motive imsys (?) tracking system which uses the **TODO:** vrpn natnet protocol to send the interaction from the 6 DOF stick in the three dimensions to the client machine. Since we are visualizing the data in two dimensions we have at least one spatial dimension that can be repurposed for secondary/other uses rather than tracking the position of the stick from world coordinates into screen coordinates.

Table 1: placeholder comparisons of different protein similarity comparisons/algorithms

dataset	full performance (fps)	half performance (ms)
big (3k images)	1,243	0.1
medium (300 image)	23	23
small (12 images)	23,312,134.3	22.1

The tracking stick also has two buttons that can be interacted with yielding a total of 4 buttons states (off off, off on, on off, on on) that can be used.

Our design process of the interactions was first laying out all of the possible actions that we wanted to cover and then experimentally incorporating them in to the implementation.

We wanted to be able to

1. have a cursor

2. zoom into the picture and explore the different clusters

3. navigate into a cluster

4. display a cluster as a gridview if a) the cluster size allowed this and b)there was enough screen real estate available for this.

5 DISCUSSION OF RESULTS

As we saw on the last few pages, we have written a plugin for MegaMol that takes pdb datafile or already generated molecular maps as an input and clusters the corresponding individual protein maps by similarity.

5.1 comparing our findings with other protein similarity practices

Laa di daa bla blubb I am a placeholder aylmao, look at me. bla [8] blah protein database simlilarity measure.

6 CONCLUSION

ACKNOWLEDGMENTS

REFERENCES

[1] S. Grottel, M. Krone, C. Müller, G. Reina, and T. Ertl. Megamola prototyping framework for particle-based visualization. *IEEE transactions on visualization and computer graphics*, 21(2):201–214, 2015.

[2] M.-K. Hu. Visual pattern recognition by moment invariants. *IRE transactions on information theory*, 8(2):179–187, 1962.

[3] T. Karrels. Fourier descriptors for shape recognition, 2006.

[4] I. Kolesár, J. Byška, J. Parulek, H. Hauser, and B. Kozlíková. Unfolding and interactive exploration of protein tunnels and their dynamics. In *Proceedings of the Eurographics Workshop on Visual Computing for Biology and Medicine*, pages 1–10. Eurographics Association, 2016.

[5] J. Krause, J. Deng, M. Stark, and L. Fei-Fei. Collecting a large-scale dataset of fine-grained cars. 2013.

[6] A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in neural information processing systems*, pages 1097–1105, 2012.

[7] M. Krone, F. Frieß, K. Scharnowski, G. Reina, S. Fademrecht, T. Kulschewski, J. Pleiss, and T. Ertl. Molecular surface maps. *IEEE transactions on visualization and computer graphics*, 23(1):701–710, 2017.

[8] D. La, J. Esquivel-Rodríguez, V. Venkatraman, B. Li, L. Sael, S. Ueng, S. Ahrendt, and D. Kihara. 3d-surfer: software for high-throughput protein surface comparison and analysis. *Bioinformatics*, 25(21):2843–2844, 2009.

[9] C. Müller, M. Krone, K. Scharnowski, G. Reina, and T. Ertl. On the utility of large high-resolution displays for comparative scientific visualisation. In *Proceedings of the 8th International Symposium on Visual Information Communication and Interaction*, pages 131–136. ACM, 2015.