

COMVI - Comparative Visualization of Molecular Surfaces using Similarity-based Clustering

Wilhelm Buchmüller, Shoma Kaiser, Damir Ravlja, Enis Spahiu

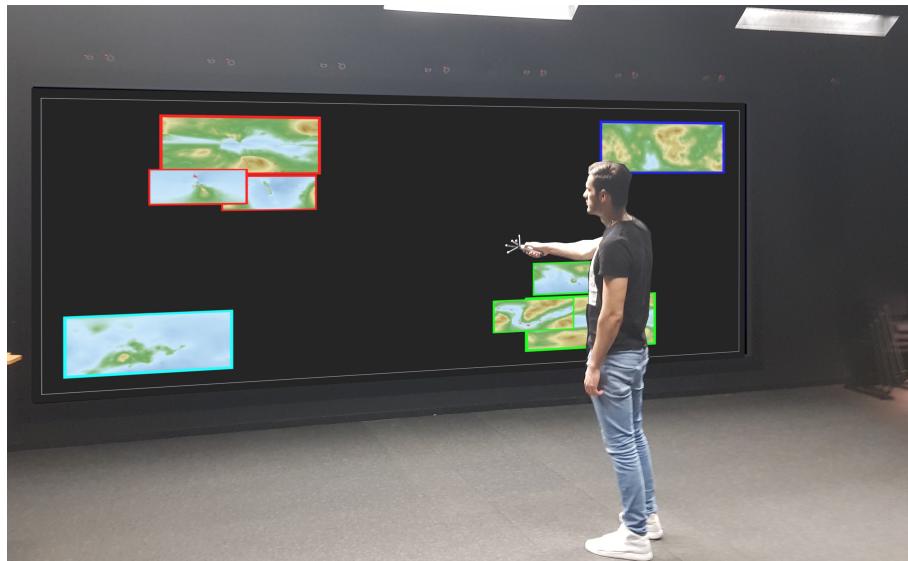


Figure 1: Mockup of a MegaMol/MSMCLUSTER instance running on the Powerwall

Abstract— The goal of this paper is to show the reader the abstract methods and concrete applications that were used to cluster proteins by similarity. We built a plugin MSMCluster for MegaMol framework that visualizes the results of clustering in a novel way. As representations of protein surfaces we used Molecular Surface Maps provided by Krone and Frieß [1]. To extract the features from Molecular Surface Maps image moments were used. Using the extracted features we clustered the data using different clustering algorithms like K-means and hierarchical agglomerative clustering. To visualize the data we used representative images of clusters in a force directed layout. To further enable the user to analyze the results of clustering, the ability to perform and visualize the clustering one level deeper was included. Since the researchers might work with large data sets, plugin was extended in order to include the possibility of interacting with large high-resolution display Powerwall. In order to test the clustering results we created three different data sets of Molecular Surface Maps. We compared the results of different clustering algorithms applied to those data sets. The paper describes the process and the approaches that were taken to solve these tasks.

Index Terms—Clustering, Dimensionality Reduction, Similarity, Feature Extraction, Visualization, Large High-Resolution Display, MegaMol, Powerwall

1 INTRODUCTION

Over the span of the last 6 months we have researched and implemented a comparative clustering process for images of molecular maps. The task consisted of several parts: to generate Molecular Surface Map images with the existing `molecularmaps` plugin for MegaMol [2], to build up a data set, to extract a descriptive feature vector from the image, to find a way to cluster the extracted feature vectors by similarity and to visualize the resulting clustering on the large high-resolution display “Powerwall”. The interaction with the Powerwall was to be done with an existing tracking device.

The results of this work are MSMCluster plugin that extends the functionality of the MegaMol visualization framework and comvi, a clustering engine that can be called by MSMCluster, or used independently. MSMCluster contains two modules that cluster

Molecular Surface Maps and visualize the results of the clustering. ClusterRenderer module visualizes the clustering using the force directed layout and in ScatterPlotRenderer individual images are represented as points in two-dimensional space. This allows the user to see the distances between images.

MegaMol is a simulation tool developed by the University of Stuttgart and the TU Dresden. It can be used to visualize particle data, simulations on atomic scale and other molecular processes such as the one that we are presenting in our work. Due to its modular nature, it can be extended with modules to either build standalone plugins or build modules that interact and interface with each other.

The first task was to retrieve molecular image data through existing MegaMol plugins [1]. For this a special binary of megamol was compiled and will be released in a separate project that will be publicly available starting May 31, 2018¹. The next task was to extract an expressive feature vector from those images and to cluster them by similarity. The last task which was also developed in a plugin in MegaMol was the visualization on the VISUS Powerwall.

Wilhelm Buchmüller
wbuch.willi@googlemail.com
Shoma Kaiser
shoma.kaiser@googlemail.com
Enis Spahiu
enis.spahiu@hotmail.de
Damir Ravlja
st144386@stud.uni-stuttgart.de

Student Project (ProjINF) started 1 December 2017; finished 31 May 2018

¹<https://github.com/aiosin/comvi>
Last retrieved: May 31, 2018

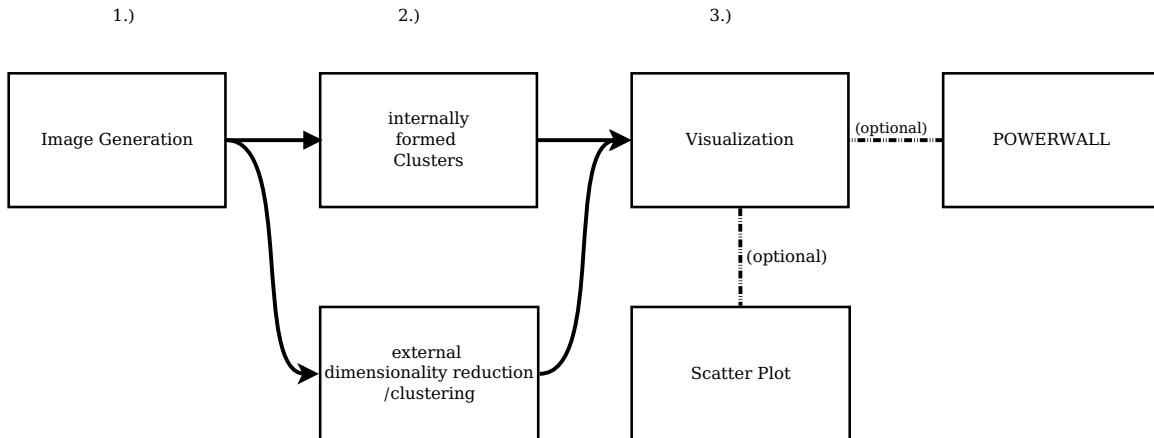


Figure 2: Architecture of the system that was implemented in comvi and MSMCluster. 1.) describes the image generation routine with the precompiled binary of MegaMol 2.) describes the (either internally in MegaMol or externally in comvi) clustering and dimensionality routine 3) describes the visualization routine with the ClusterRenderer and optionally with the ScatterPlotRenderer or on the Powerwall

The Powerwall is a very high resolution display² that can be used to visualize large data(sets). Due to its size and resolution it is possible to display much more information than on a regular consumer grade screen.

The Powerwall supports a tracker device that can transmit 6 degrees of freedom, so for the interaction part we had more possibilities than with traditional human interface devices (HIDs). For this last step we also researched the possible interactions with this unique tracking device which can be used to interact with the Powerwall.

The task required from us that we learn how to compare the images, measure the similarity distances between the images, and cluster them. Furthermore, it required that we use a similarity based clustering algorithms.

To prevent confusion further down the paper the reader should be familiar with the following two terms: `comvi` - internal name of the clustering engine that can be used by the MegaMol Plugin `MSMCluster` - name of the Megamol plugin that was developed as the main part of the project.

Over the course of the next few pages you will learn how we approached these challenges and how we solved them. You will find out what worked and what did not. In the second chapter related work is presented. Section 3.5 is the main part of the paper and shows our approaches to clustering and extracting features from the protein images that were used in the clustering. In section 4 our approaches to the visualization of data on conventional desktop computers and large high-resolution displays are presented. In this section we also describe shortly the visualization with the `MSMCluster` plugin. In section 5 we present and discuss the results of applying `MSMCluster` plugin onto data sets of different sizes. In the conclusion we shortly discuss possible future work.

2 RELATED WORK

Clustering proteins by similarity or comparing individual proteins has been subject of existing work.

The paper “Unfolding an interactive exploration of protein tunnels and their dynamics” [3] from Kolesar et al. already has similar approaches to ours. Kolesar et al. used a 10 dimensional feature vector based on invariant image moments defined by Hu [4]. Hu-Moments are set of invariant moments defined over a two dimensional signal, that give unique features about the object in the image. More about image moments and the Hu Moments in section 3.3.1.

²<https://www.visus.uni-stuttgart.de/institut/visualisierungslabor/technischer-aufbau.html>
Last retrieved: May 31, 2018

Another approach of 3D protein data were 3D Zernike moments explored by La et al. in [5]. The approach is fundamentally the same as in Kolesar et al. but La et al. use Zernike moments instead of traditional image moments and extended them to three dimensions.

For further information and more complex approaches to this problem one should consult the Smith Waterman algorithm from Smith and Waterman “Identification of Common Molecular Subsequences”[6]. This work considers the protein as a one-dimensional sequence of amino acids, rather than extracting two or three dimensional features over a two or three dimensional signal.³ A detailed description would go beyond the scope of the paper, but in short the algorithm calculates the optimal alignment of two sequences, for which a local similarity matrix of the two sequences is computed. The matrix is then used to backtrack the optimal global alignment from a local optimum.

Visualization of data on large high-resolution displays is discussed in Müller et al. [7]. The System that they built could be used on large high-resolution displays as well as consumer grade desktops. It allowed comparison of simulation results. They also conducted a user study where the participants could compare different views of the protein in the simulation. They had to choose the view that was closest to the baseline. Their results showed that, depending on the data set, users were either more or equally precise and confident that they found the correct solution when using large high-resolution displays, compared to when they used visualization on conventional displays. The System used for the user study was showed to the experts in the field of technical biochemistry and they said that this way of visualizing the data would possibly allow new discoveries in some areas of their research.

3 CLUSTERING

Due to the given generated data not being labeled, only unsupervised machine learning techniques like clustering could be applied on it. Clustering of data is defined as the process of grouping n distinct values into $m \leq n$ different classes. Clustering algorithms are algorithms that perform clustering. They can be subdivided in two different ways [8].

Depending on the number of classes to which one data point can belong, one can differentiate between hard and soft clustering. In hard clustering one element can belong to only one class, whereas in soft clustering it can belong to many different classes. On the other hand, depending on the way in which the data is clustered

³An overview can be found here: <https://www.ebi.ac.uk/Tools/ssss/> Last retrieved May 31, 2018

one can differentiate between flat and hierarchical clustering. In flat clustering particular clusters have no relation to each other, whereas in hierarchical clustering one considers the distance between clusters as well, this results in a hierarchy.

3.1 Approaches to the Clustering-Problem

At the start of the project we had several ideas of how we could approach this problem. Trying to follow the recent trend in machine learning we researched possible ways to incorporate neural networks to compute a similarity metric between two images or classify an image into a more usable vector of data.

We ended up using a higher dimensional feature vector described in section 3.2 to compute the similarity between two protein maps because we did not manage to train a custom neural network model in the given timeframe, partly due to lack of knowledge in the field of machine learning and also due to the lack of labeled data, which would have been required.

If one had labeled data, one could train an autoencoder, SVM or simple neural net (if the number of classes is relatively low).

Our relatively spartan results with a pretrained Imagenet [9]⁴ model led us to believe that given the knowledge on the subject, and data labeled by human experts in the field, it should be possible for this specialized task to find a machine learning solution using neural networks/autoencoders.

3.2 Finding a feature vector to cluster the images

This section describes the process of finding a feature vector that describes the Molecular Surface Maps. Some of those feature vectors were then used as an input to the clustering algorithms used by the MSMCluster plugin (see Figure 2).

The challenge of finding a good feature vector is to find good features which can give a lot of information about the image. If the extracted image features do not describe the images in a useful way, dimensionality reduction will give noisy results and the clustering algorithms will not group the images in a meaningful way, so users will not be able to extract any new information from the clustering.

After finding the feature vector for a given image, the following procedure is to apply some sort of dimensionality reduction to project a higher dimensional vector onto a 2D or 3D space.

This has multiple advantages one of which is that if the dimensionality reduction works as intended one can easily find out after applying the dimensionality reduction if similar looking items are positioned next to each other.

3.2.1 Curse of dimensionality

Alternatively one could also find a similarity measure for higher dimensional vectors, but this problem has turned out harder than it seemed partly due to the *curse of dimensionality*⁵. Contrary to expectations, in very high or infinite dimensional spaces, unexpected things start to happen. If we such as the euclidean distance or mathematically put the L_2 -norm loses relevance since:

$$\lim_{n \rightarrow \infty} x^n = 0 \quad \forall x \in [0, 1] \quad (1)$$

Simply put, otherwise very dissimilar values have their distance drawn to zero. For example between a vector with 0.99 in every dimension and a vector with 0 in every dimension the euclidean distance is: $0.99^{100} = 0.33660$ which would put them closer than one would assume them to be.

So we have to come up with other solutions to this problem discussed in section 3.4.

After looking at a small subset of molecular map images from a variety of proteins we assumed that we could extract feature information about the following properties of a given image: Color distribution, Shape, Texture and Invariant Image Moments.

⁴<https://pjreddie.com/darknet/imagenet/>
Last retrieved: May 31, 2018

The initial idea was that the color distribution gives information about the color palette and local color changes in the image, the extracted shape features should give information about the the biggest n shapes in the picture, the texture feature should differentiate between smooth, rough and coarse texture and everything in between.

The image moments were chosen as a go to approach to extract invariant features from the image which has been proven to yield usable results as described in Kolesar et al. [3].

3.3 Image Features/Descriptors

This section explains the different image features that we used to build the feature vector. For the length of this section the reader should know that in `comvi` we rescale all images to the size of (128, 128) after reading them into our routine, whereas in `MSMCluster` image moment calculation is applied on the original image.

3.3.1 Invariant Image Moments by Hu

The set of invariant Image Moments discovered by Hu et. al are rotation, translation, scale and transformation invariant. This allows us for example to determine if an image I_a is similar to another image I_B if I_B is equal to I_A and simply rotated by 30°.

The continuous image, defined over the domain \mathbb{R}^2 over two dimensions are defined as such:

$$M_{p,q} = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} x^p y^q f(x,y) dy dx$$

Where $f(x,y)$ is the signal of the image.

When dealing with discrete values in a discretized RGB or greyscale (GS) image, we use sums instead of integrals, so we get the following formula:

$$M_{i,j} = \sum_{x=0}^n \sum_{y=0}^m x^i y^j f(x,y)$$

Since Kolesar et al. found that a set of 10 image moments invariant to translation was sufficient to describe the differences between grayscale images of protein surfaces, image representation in `MSMCluster` was done using their set of image moments. However, since Molecular Surface Maps are color images, another image representation was included in `MSMCluster`. This image representation used three ten-dimensional sets of image moments invariant to translation that were calculated for every RGB color channel. In `comvi` on the other hand we used the standard Hu et al. invariant moments.

3.3.2 Color palette/histogram

The goal when extracting color palette was to reduce the big color space that is present in any of the molecular maps and get a few distinct colors that describe the color space with a few representative colors.

To achieve this we extract a histogram of each color channel of the RGB images. Each channel is then represented as a greyscale image. We then create for each channel a histogram of the luminance intensities with, in our case 16 bins. This “reduces” the $128 * 128 = 16.384$ dimensions to just 48 dimensions for our image.

An alternative approach that we used to extract color information about the image is to split the image in to equally sized chunks and to extract simple color distribution features like average, min and max values.

The results from both approaches will be discussed in 3.4.

3.3.3 Texture

After looking at the smallest data set of molecular maps we noticed that many samples had a distinct roughness that looked like they could be used to classify their texture.

We ended up using the Haralick textural features [10]. The Haralick features work with a grey level correlation matrix (GCM). The GCM for a given greyscale image is defined as follows:

$$C_{\Delta x, \Delta y}(i, j) = \sum_{x=1}^n \sum_{y=1}^m \begin{cases} 1, & \text{if } I(x, y) = i \text{ and } I(x + \Delta x, y + \Delta y) = j \\ 0, & \text{otherwise} \end{cases}$$

The Haralick features compute, based on this matrix, features that aim to describe the texture of the original source image.

The features include e.g. the angular second (image) Moment, entropy, summed variance/energy/average.⁶

Also we can take more features from that GCM which gives us more features to work with.⁷

We also computed this on every channel yielding us 3×14 features.

Alternatively one can also use the texture features developed by Tamura et al. in “Textural Features Corresponding to Visual Perception”[11]. The Tamura features are different from the Haralick features since they were developed on a basis of psychological experiments. The Tamura features do not use a GCM, instead they use a combination of kernel convolutions that yield the information about the coarseness, contrast, directionality, linelikeness, regularity and roughness.

3.3.4 Shape

The shape features we researched were fourier descriptors^{8 9}. In short, we get the centroid-distance curve of a distinct shape and take the fast (discrete) fourier transform of that curve. The technique has been proven to work for shape recognition in earlier works [12].

Our approach was the following. Since our molecular maps are extremely complex (image-wise) we first need to divide the image into discrete image regions. We preferably want to segment our image into n colors. To recap: our everyday monitor does support 24 bit colors, 8 bit per channel. This gives us $2^{24} = 16.777.216$ total colors. To find a valid contour in that colorspace would be a coincidence in most cases. Our approach is therefore to compress the color space down to a couple of colors. This proces is called color segmentation. We use the K-means algorithm to put every color from our original image into a few bins.

Alternatively the mean-shift algorithm can be used. The algorithm is a kernel based clustering algorithm which operates sort of like the gradient descent algorithm. The mean shift algorithm is explained in detail in section 3.5.

A comparison of color segmentation of mean shift and k-means can be seen here^{10 11}.

Nevertheless our problem of color segmentation was still not solved. Especially in the k-means algorithm we encountered a lot of smaller “irrelevant” shape artifacts that we did not want. So one way of getting them out of the way was to squash the color space further down to two or three distinct colors in the image.

The other way is to dynamically (depending on each image) dilate and erode¹² the image until we can detect with a matrix labeling or

⁶http://murphylab.web.cmu.edu/publications/boland/boland_node26.html

Last retrieved: May 31, 2018

⁷http://murphylab.web.cmu.edu/publications/boland/boland_node26.html

Last retrieved: May 31, 2018

⁸<http://fourier.eng.hmc.edu/e161/lectures/fd/node1.html>

⁹<http://demonstrations.wolfram.com/FourierDescriptors/>

¹⁰<https://spin.atomicobject.com/2015/05/26/mean-shift-clustering/>

¹¹<https://spin.atomicobject.com/2016/12/07/pixels-and-palettes-extracting-color-palettes-from-images/>

Last retrieved May 31, 2018

¹²https://docs.opencv.org/2.4/doc/tutorials/imgproc/erosion_dilatation/erosion_dilatation.html

Last retrieved May 31, 2018

floodfill algorithm that we really have only a handful of shapes in the image.

For the shapes we then extract the contours and compute the centroid of that shape. The centroid is the geometric center of mass of the image which can be built up by the first central moments of the two dimensions.

Now we have the centroid and the contour of the shape that we want to classify. We need to form the contour centroid distance curve. The contour centroid distance curve is simply a discrete list of values of the x and y values of each item of the contour or alternatively the list of the euclidian distances from the centroid. Both approaches have been shown to work.

We then normalize the values by diving all the values by the biggest value in the array, resulting in an array that is in range $(0, 1]$.

Without loss of generality depending on which curve we chose we take the fast discrete fourier transform of that signal, which gives us for each sine/cosine coefficient a weight of how much that component contributes/weights in the signal.

Without any proof that would go beyond the scope of this paper we can say that the gained features are rotation scale and translation invariant. It is rotation invariant because the Fourier transform is shift invariant so it does not matter where we begin our centroid curve; the power spectra of the Fourier transform will be the same even if its rotated. It is scale invariant because we normalize it and it is translation invariant because the distance curve itself is translation invariant.

Another shape feature can be won, under the condition that the shape object is given as a greyscale binary matrix. From this matrix one can extract Image Moment features like the one by Hu as described in section 3.3.1.

We did not end up using the shape features, because we could not notice any improvement over the other features that we used, but we put substantial effort into getting these features to work, so we wanted to describe our process here.

3.4 Finding the best performing dimensionality reduction method

For testing purposes we used every dimensionality reduction method we could find that looked halfway promising. We tested on data sets of various sizes and content.

We tested our results on the medium sized Oxford flower data set[13], the large sized car data set [14] and our large (~ 3600 images) molecular maps data set. We further extracted a smaller sized (~ 300 images) data set as our benchmark set and an easy to benchmark and a small (12 images) data set for faster testing.

As you will find in our `comvi`repository¹³ the dimensionality reduction routine consists of passing it a feature array of shape $A^{n \times m}$ where n is the number of samples (i.e. the number of images) and m is the number of features per sample. In the routine we additionally scale the data because we want to avoid potentially unwanted effects caused by very big feature values.¹⁴ We then return the reduced feature array back to the callee. The callee can then apply some further functions on this reduced array.

After initial testing with linear dimensionality reductions we notice that our results just would not “converge”, so we again tried with any non-linear dimensionality reduction technique we could find.

We tested our data with Local Linear Embedding (LLE), Multi-Dimensional Scaling (MDS), Kernel-PCA, Isomap and t-distributed Stochastic Neighbor Embedding (t-SNE). After testing with a small non representative user group that would judge how good the clustering was done by the technique we concluded that t-SNE performed the best.

For a early test with linear dimensionality reduction rather than non-linear/manifold-learning techniques one should look at Figure

¹³<https://github.com/aiosin/comvi>

¹⁴<https://medium.com/greyatom/why-how-and-when-to-scale-your-features-4b30ab09db5e>

Last retrieved: May 31, 2018

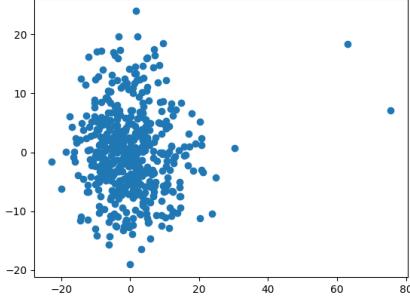


Figure 3: Early test with linear PCA on the BMW car data set [14]

3, where a linear dimensionality reduction technique was applied, opposed to Figure 5 and Figure 4 where the t-SNE algorithm was applied.

3.5 Clustering Algorithms

We applied the following clustering algorithms to image descriptors that were extracted from the protein images.

K-means takes the number of clusters as an input and if the data set contains more elements than the wanted number of clusters, groups elements into that many clusters. Since it always groups data into the specified number of clusters, K-means is a flat clustering algorithm. Because every image is clustered into only one cluster it belongs to hard clustering algorithms.

To cluster the data K-means uses the concept of centroids. A centroid is a point in vector space that is located in the middle of the corresponding cluster. There is therefore one centroid for every cluster. At first K-means initializes the same number of centroids the same as the number of clusters. This can be done by choosing randomly k samples in the vector space. Since K-means is an iterative algorithm it then repeats reassignment and recomputation steps until it minimizes the distance between the centroid and all of the elements of the corresponding cluster as follows [15]:

Let $D = \{x_i\}_{i=1}^n$ be a set of n vectors x_i that represents some data. This data should be clustered into k clusters whose centroids are $\mu_{c(i)}$. Now let $c : i \mapsto k$ be the assignment of the element i to the cluster k . K-means minimizes the following loss function L :

$$L = \min_{c, \mu} \sum_{i=1}^n (x_i - \mu_{c(i)})^2 \quad (2)$$

This minimization is done by repeating the reassignment (3) and recomputation (4) step.

$$\forall i : c(i) := \arg \min_k (x_i - \mu_k)^2 \quad (3)$$

$$\forall k : \mu_k := \frac{1}{|c^{-1}(k)|} \sum_{i \in c^{-1}(k)} x_i \quad (4)$$

K-means does not converge to a global minimum which is one of its main weaknesses. In order to circumvent this problem centroids are usually initialized randomly, but this makes the algorithm non-deterministic. To find the best possible clustering several restarts of the algorithm are sometimes needed. Another disadvantage is that K-means finds only spherical clusters because it adds an image only to the cluster whose corresponding centroid is nearest to that cluster [16]. The

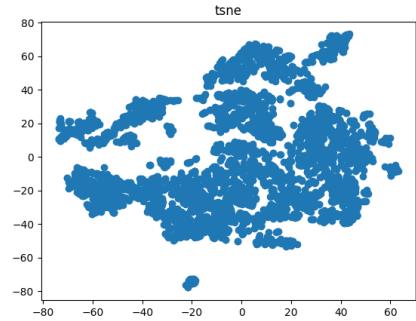


Figure 4: Test with the tsne dimensionality reduction on our large (~ 3600 images) dataset

main advantages of K-means on the other hand are that it often gives good results and is easy to implement. It is also efficient because K-means time complexity is linear in the number of elements, number of clusters and number of iterations.

The `MSMCluster` plugin contains two implementations of K-means. We implemented one version ourselves and the other version is part of the `dlib` library [17].

Hierarchical agglomerative clustering is a type of hierarchical clustering that clusters the data from the bottom up. As an input it takes the wanted number of clusters. It starts by making a cluster for every element in the data set. Then repeatedly merges clusters that are the most similar according to some metric until in the end there is only one cluster left. However in the course of every merging operation it takes a note of the similarity between the clusters that are merged. This way it builds cluster hierarchy that can be visualized with dendograms¹⁵. The given number of clusters is then obtained by cutting the tree on the right similarity level. This is done by descending from top of the tree and taking into account the number of siblings on every similarity level [8].

Main advantages of hierarchical agglomerative clustering is the hierarchy of clusters that it creates, and its stability over multiple executions of clustering. Its main weakness is on the other hand the high time ($O(n^2 \log(n))$) and space complexity ($O(n^2)$) [18].

The `MSMCluster` plugin uses implementation of hierarchical agglomerative clustering that is contained in the `dlib` library.

Mean-shift is another clustering algorithm that relies on the concept of centroids. It is a density based algorithm much like DBSCAN [17]. Every point which represents a sample in two dimensions spans a circle around it with a preconfigured bandwidth size (this size can be algorithmically determined). In a single iteration step of the algorithm every sample is moved towards the centroid of all the points that the kernel of the sample contains. This step is repeated until every cluster has reached a local minimum and no more points can be accommodated in that cluster. The bigger the bandwidth of the kernel the smaller the amount of clusters.¹⁶. For a demonstration of the results produced by the mean shift algorithm look at Figure 5.

¹⁵A dendrogram (from Greek *dendro* "tree" and *gramma* "drawing") is a tree diagram

¹⁶<https://spin.atomicobject.com/2015/05/26/mean-shift-clustering/>
Last retrieved: May 31, 2018

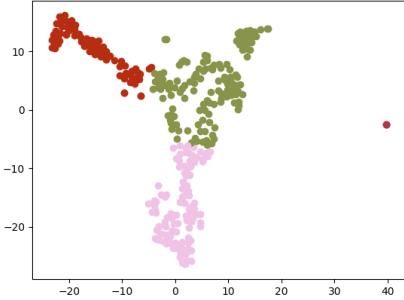


Figure 5: Final test with the t-SNE dimensionality reduction and the mean shift clustering algorithm

3.6 Finding the best performing clustering algorithm

Our approach to finding a clustering algorihtm was the same as the one that we used to find the best performing dimensionality reduction algorithm.

Except with a little modification. Kolesar et al. [3] use the k nearest neighbors algorithm over their 10-dimensional feature vector to cluster their data, as they discuss their results. We decided to further use this approach to cluster our molecular maps instead of the molecular protein tunnels which Kolesar et al. analyzed.

In the early stages of the project we intended to use the DBSCAN clustering algorithm. The key feature to this algorithm is that it automatically determines the optimal number of clusters that it detects in the given data.

The algorithm is a density based spatial clustering algorithm that uses a similarity metric between two items over the domain $[0, 1]$. So the ways of computing the similarity is given more freedom with this approach.

One could use the euclidian distance between two feature vectors, the cosine distance or feed the two images to a siamese neural network/autoencoder. Siamese neural network are, for the sake of simplicity, two identical black box oracle machines which yield different valued vectors if you feed the two machines two different signals. An autoencoder is a machine which will have learned to classify data into a few distinc categories. The main disadvantages of these techniques is that both the siamese neural network and the autoencoder have to be trained.

One could also use the Haussdorff distance to compare two higher dimensional vectors. A detailed explanation would go beyond of the scope of this paper, but the Haussdorff distance is simply put the “maximum distance of a set to the nearest point in the other set”¹⁷.

We opted against DBSCAN for our primary clustering algorithm because in the early stages of the project the computation of the $n \times n$ similarity matrix was computationally expensive and took more time than the approaches that we demonstrate elsewhere.

For our initial testing we used k-nearest neighbors with a preconfigured assumed number of clusters.

We further tried a spectral clustering ¹⁸ and chinese whispers clustering algorithm[19]. Both of these algorithms are graph based algorithms. The spectral clustering approach works with the eigenvalues (spectrum) of the similarity matrix (like DBSCAN), while the chinese whispers clustering does reachability analysis in a network of connected nodes. Both of these approaches yielded unusable results, since the chinese whispers algorithm for example is not suitable for this task and performs poorly on any similarity matrix.

¹⁷<http://cgm.cs.mcgill.ca/~godfried/teaching/cg-projects/98/normand/main.html>

¹⁸https://www.cs.cmu.edu/~aarti/Class/10701/readings/Luxburg06_TR.pdf

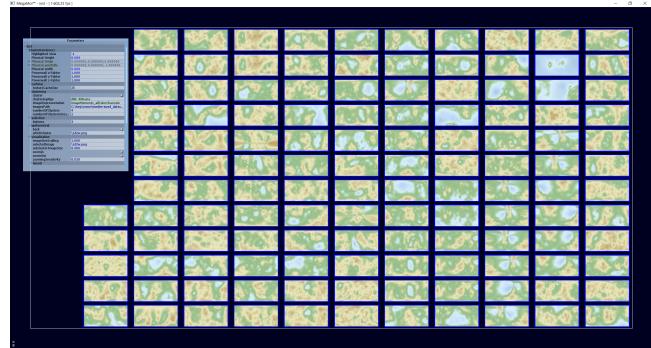


Figure 6: Visualization of all images within one cluster in a uniform grid layout. In the top left corner a tweak bar that contains configurable parameters can be seen.

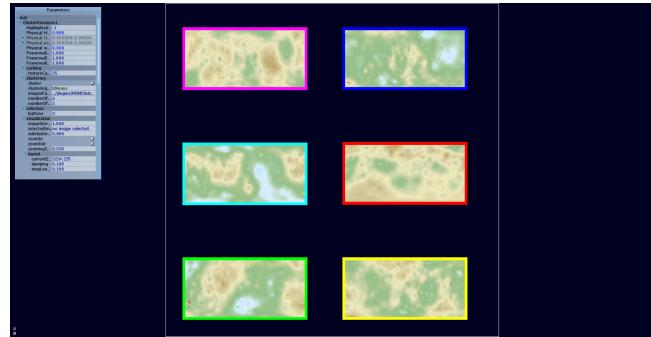


Figure 7: A running MegaMol instance, in the normal clustering view.

We ended up with the algorithms described in section 3.5. It also is a density based algorithm and much like the DBSCAN algorithm is also determines the optimal number of clusters for the given parameters but it operates with a kernel sampled on each point in the plot and calculates the center of mass so to speak of the points clouds

4 VISUALIZATION AND INTERACTION

The content of a running instance of MegaMol consists of a bounding box in which the content is drawn and a tweak bar that contains configurable parameters in named drop-down menus (see Figure 6 and Figure 7). On conventional desktop monitors the user interacts with the MSMCluster plugin using a mouse and a keyboard.

4.1 Background MegaMol and Powerwall

The visualization of MegaMol can be displayed on the Powerwall. The Powerwall is a merged large high-resolution display made of five portrait-oriented strips. The setup has a resolution of 10.800 × 4.096 pixels for each eye and is projected onto a physical screen size of about 6 × 2.2 metres. Two 4k (Liquid Crystal on Silicon) LCoS projectors are projecting the images for each strip, for a total of five strips. A pixel is about 0.56 mm in size, which corresponds to a density of about 50 ppi. The user interacts with the system using a 6-DOF mid-air pointing Stick with two buttons (1 on each side), which is tracked by 14 cameras using an optical tracking system.

4.2 Realization

After clustering the Molecular Surface Maps, the MSMCluster plugin then visualizes the results of clustering. Both modules contained in the plugin deal with visualization of the clustering.

The ClusterRenderer module first extracts from every cluster an image that represents that cluster. If the clustering algorithm uses the concept of centroids, it chooses the nearest image to the centroid as a representative image of every cluster. If there is no such concept as a centroid, the ClusterRenderer randomly chooses an image from the cluster to represent it. Representative images are then drawn, with borders in different colors, onto the screen within the bounding box in MegaMol. To make the representative images fit better in the bounding box, the aspect ratio of the bounding box was set to 2:1, the same aspect ratio as images.

To position the images on the screen we used techniques from graph visualization. In the ClusterRenderer every representative image is considered to be a node within an imaginary graph. To determine the positions of representative images the implementation of force directed graph drawing algorithm contained in MegaMol was used. Force directed algorithms are based on the physical laws and try to imitate the effect of attractive and repulsive forces on the graph nodes, as if they were connected by springs. An implementation contained in MegaMol uses repulsive forces based on Coulomb's law and attractive forces based on Hooke's law. It also uses the concept of maximum internal energy. When internal energy becomes higher than the maximum internal energy movement of images stops. Further parameters that could be given to the force directed layout instance are damping that affects how fast the system calms down and integration step that defines the movement speed of the nodes. These parameters related to the layout are configurable in tweak bar in drop-down menu layout. In this menu the current internal energy of the top clustering level is shown to the user as well. Other parameters that are configurable in ClusterRenderer related to visualization are the size of the images for both levels of clustering. In the parameter tweak bar the currently selected image is shown as well.

ClusterRenderer also supports (1) further clustering of the obtained clustering results (see Figure 1, smaller images around larger images) and (2) showing all of the images within one cluster using grid layout. Clicking with the right mouse button on the representative image of some cluster performs clustering only on the images contained in the clicked cluster. Result of every such deeper level of clustering is visualized using the representative images of internal clusters. Those images are smaller than the image of the parent cluster and have the same border color as the parent cluster. They are positioned using another instance of force directed layout. This instance of force directed layout considers only repulsive and attractive forces between the image of the parent cluster and the images of the subclusters. If the user however clicks on the image using the left mouse button while holding the CTRL key, all of the images within the selected cluster are shown in a uniform grid layout (see Figure 6). After testing we decided that one level of subclustering is enough. After 2 levels of subclustering the results get too noisy and unusable.

To interact with the Application intuitively on PC and on the Powerwall it was important that we use movements for the interactions from our everyday life. If the Application is used on PC, the MegaMol Project is mainly interacting with the mouse. If we run Application at the Powerwall it is interacting with the tracking stick.

With the four button states of the tracking sticks we focused on the most important interactions like "selection" or "drag and drop". To localize the aiming of the user we calculate intersection coordinates of the tracking stick in direction to the Powerwall. While moving the tracking stick for interactions, these coordinates should be scaled for user-friendly interactions.

The interactions are executed by mouse or by the tracking stick. The tracking stick has got as already mentioned four button states. In the first state, if no button is pressed, there should be no interaction. In the second state, if the first button is pressed, there is a 'drag and drop'-interaction. When only pressing the other button, the third state, the user can select an image. If both buttons are pressed, the user is able to zoom. For the zooming we used the third dimension of the tracking area in front of the Powerwall. When walking towards

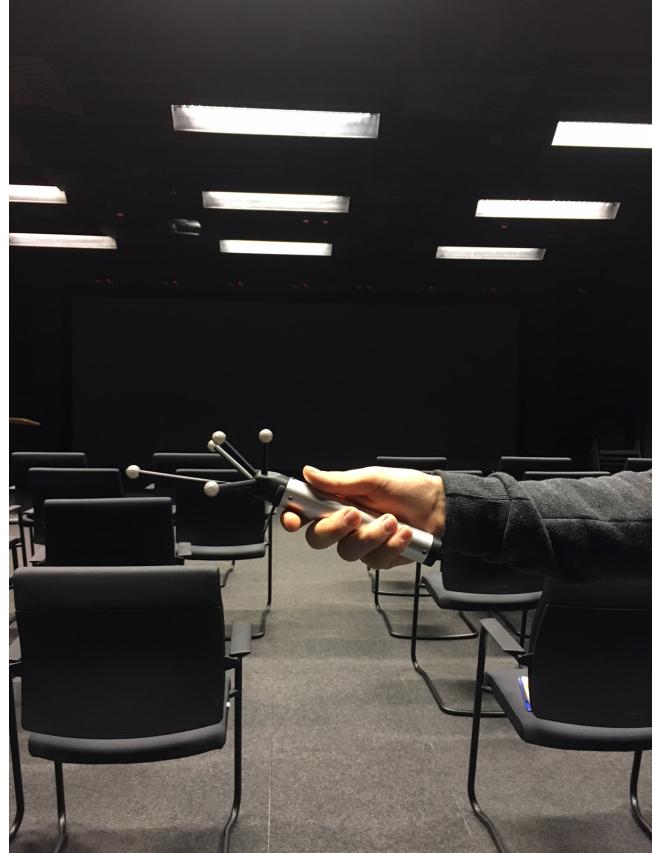


Figure 8: The tracking stick used for the interactions with the powerwall.

or away from the Powerwall in the fourth button state (when the user holds both buttons) and zooms the view. A picture of the tracking stick can be seen in Figure 8.

5 RESULTS

As we mentioned on the last few pages, we have written a plugin for MegaMol that takes a directory with generated Molecular Surface Maps as an input and clusters them by similarity. In this section we present and discuss the results of testing three different data sets with MSMClusterplugin.

Data sets were ordered by size and were used for different purposes. The smallest data set contained only 12 images and could therefore be used to test changes in the program quickly on all kinds of hardware. The second data set was a medium-sized data set that contained 306 images. This data set was used to test the plugin and help us notice the areas where the system needed to be optimized. The largest data set contained all of the 3646 protein images that we generated. This data set will be published on Kaggle¹⁹.

During the creation of Molecular Surface Maps, some of the images were not created correctly and ended up having large black areas. Those images appeared as outlying data points in scatter plots. We opted against the removal of the outliers because they were often clustered into one cluster and during plugin development this proved to be a good way to test whether the clustering algorithm is working correctly.

Since there were no labeled data sets we could not use external criteria like purity or rand index to evaluate the results of our clustering. Instead we looked at the results of clusterings and noticed

¹⁹<https://www.kaggle.com/zython/molecular-protein-maps>

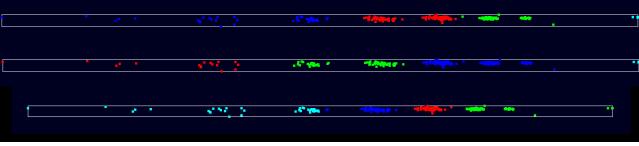


Figure 9: Visualization of the clustering results using the Scatter-PlotRenderer module when applying the following algorithms to the medium-sized data set (from top to bottom): (1) Hierarchical Agglomerative Clustering, (2) dlib implementation of K-means, (3) our implementation of K-means. Wanted number of clusters set was set to 4.

that the clustering results of different clustering algorithms produced similar results. See Figure (9).

Using different image representations however impacted the results more than using different clustering algorithms. When using Image moments computed on all color channels (see section 3.3.1) the results of clusterings looked somewhat less noisy as expected. This improvement comes nevertheless with more computational overhead i.e. since we are in this case computing image moments on all channels the, it costs three times as much time as computing on one channel (greyscale moments).

Additionally the task required that we compare our findings of our 2D clustering to existing 3D protein similarity methods. We used the 3D-Surfer application for our comparison [5]. However we noticed that, because even our largest datasets are significantly smaller than the datasets used by 3D-Surfer ($3646 \ll \sim 600.000$), we could not compare the two methods. A lot of proteins that were given by the 3D Surfer application were not available in our dataset and to build up the complete dataset of 600.000 images would have taken up a significant portion of our time budget.

6 CONCLUSION

In this paper we looked at different ways of clustering proteins by similarity, visualizing results of the clustering and allowing the user to interact with the data on wall-sized high-resolution display Powerwall. The result of our research was the MegaMol Plugin MSMCluster and the clustering engine comvi.MSMCluster allows the user to cluster protein images and visualize them. This plugin can be used with the large high-dimension display Powerwall.

These Proteins were represented by Molecular Surface Maps that were extracted using the molecularmaps plugin. We found that clustering results obtained with different clustering algorithms did not differ significantly, because of that we assume that further improvements in clustering of protein images might be achieved in the research of features that are extracted from the images. Perhaps further adaption of features specifically to the Molecular Surface Maps would offer further improvements of the results.

Furthermore, with the continuation of the Moore's Law large high-resolution displays might become even more common. Because of that, further research in this area and adaptation of scientific visualization tools and other programs to these displays might play a significant role in the future.

ACKNOWLEDGMENTS

We would like to thank our supervisors Michael Krone and Florian Frieß as well as our project examiner Prof. Ertl, for giving us this opportunity to work on this project. We are grateful that we were able to improve our knowledge and learn new things. We are also grateful for the feedback we received on our work.

REFERENCES

- [1] Michael Krone, Florian Frieß, Katrin Scharnowski, Guido Reina, Silvia Fademrecht, Tobias Kulschewski, Jürgen Pleiss, and Thomas Ertl. Molecular surface maps. *IEEE transactions on visualization and computer graphics*, 23(1):701–710, 2017.
- [2] Sebastian Grottel, Michael Krone, Christoph Müller, Guido Reina, and Thomas Ertl. Megamol — a prototyping framework for particle-based visualization. *IEEE transactions on visualization and computer graphics*, 21(2):201–214, 2015.
- [3] Ivan Kolesár, Jan Byška, Julius Parulek, Helwig Hauser, and Barbora Kozlíková. Unfolding and interactive exploration of protein tunnels and their dynamics. In *Proceedings of the Eurographics Workshop on Visual Computing for Biology and Medicine*, pages 1–10. Eurographics Association, 2016.
- [4] Ming-Kuei Hu. Visual pattern recognition by moment invariants. *IRE transactions on information theory*, 8(2):179–187, 1962.
- [5] David La, Juan Esquivel-Rodríguez, Vishwesh Venkatraman, Bin Li, Lee Sael, Stephen Ueng, Steven Ahrendt, and Daisuke Kihara. 3d-surfer: software for high-throughput protein surface comparison and analysis. *Bioinformatics*, 25(21):2843–2844, 2009.
- [6] T.F. Smith and M.S. Waterman. Identification of common molecular subsequences. *Journal of Molecular Biology*, 147(1):195 – 197, 1981.
- [7] Christoph Müller, Michael Krone, Katrin Scharnowski, Guido Reina, and Thomas Ertl. On the utility of large high-resolution displays for comparative scientific visualisation. In *Proceedings of the 8th International Symposium on Visual Information Communication and Interaction*, pages 131–136. ACM, 2015.
- [8] Hinrich Schütze, Christopher D Manning, and Prabhakar Raghavan. *Introduction to information retrieval*, volume 39. Cambridge University Press, 2008.
- [9] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in neural information processing systems*, pages 1097–1105, 2012.
- [10] R. M. Haralick, K. Shanmugam, and I. Dinstein. Textural features for image classification. *IEEE Transactions on Systems, Man, and Cybernetics*, SMC-3(6):610–621, Nov 1973.
- [11] H. Tamura, S. Mori, and T. Yamawaki. Textural features corresponding to visual perception. *IEEE Transactions on Systems, Man, and Cybernetics*, 8(6):460–473, June 1978.
- [12] Tyler Karrels. Fourier descriptors for shape recognition. https://homepages.cae.wisc.edu/~ece533/project/f06/karrels_ppt.pdf, 2006.
- [13] M-E. Nilsback and A. Zisserman. A visual vocabulary for flower classification. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, volume 2, pages 1447–1454, 2006.
- [14] Jonathan Krause, Jia Deng, Michael Stark, and Li Fei-Fei. Collecting a large-scale dataset of fine-grained cars. 2013.
- [15] Daniel Hennes. Grundlagen der künstlichen Intelligenz, lecture notes. <https://ipvs.informatik.uni-stuttgart.de/mlr/teaching/ki2017/>, January 2018.
- [16] Anil K Jain. Data clustering: 50 years beyond k-means. *Pattern recognition letters*, 31(8):651–666, 2010.
- [17] Davis E. King. Dlib-ml: A machine learning toolkit. *Journal of Machine Learning Research*, 10:1755–1758, 2009.
- [18] Anil K Jain, M Narasimha Murty, and Patrick J Flynn. Data clustering: a review. *ACM computing surveys (CSUR)*, 31(3):264–323, 1999.
- [19] Chris Biemann. Chinese whispers: an efficient graph clustering algorithm and its application to natural language processing problems. In *Proceedings of the first workshop on graph based methods for natural language processing*, pages 73–80. Association for Computational Linguistics, 2006.